

Coding Lab: Manipulating data with dplyr

Ari Anisfeld

Summer 2020

```
## Warning: package 'tidyverse' was built under R version 3.6.3
## Warning: package 'ggplot2' was built under R version 3.6.3
## Warning: package 'tibble' was built under R version 3.6.3
## Warning: package 'tidyr' was built under R version 3.6.3
## Warning: package 'readr' was built under R version 3.6.3
## Warning: package 'purrr' was built under R version 3.6.3
## Warning: package 'dplyr' was built under R version 3.6.3
## Warning: package 'stringr' was built under R version 3.6.3
## Warning: package 'forcats' was built under R version 3.6.3
## Warning: package 'readxl' was built under R version 3.6.3
```

Data manipulation with dplyr

Once you have data in R, you'll want to explore it.

The tidyverse package dplyr provides a toolkit for data manipulation.

We will cover:

- ▶ `select()` to pick columns
- ▶ `arrange()` to order the data
- ▶ `mutate()` to create new columns
- ▶ `filter()` to get rows that meet a criteria
- ▶ `summarize()` to summarize data

selecting columns with select()

select()

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



storm	pressure
Alberto	1007
Alex	1009
Allison	1005
Ana	1013
Arlene	1010
Arthur	1010

selecting columns with select()

Use case: You want to present a subset of your columns

```
select(texas_housing_data, city, date, sales, listings)
```

```
## # A tibble: 8,602 x 4
##   city      date sales listings
##   <chr>   <dbl> <dbl>    <dbl>
## 1 Abilene 2000      72      701
## 2 Abilene 2000.     98      746
## 3 Abilene 2000.    130      784
## 4 Abilene 2000.     98      785
## 5 Abilene 2000.    141      794
## 6 Abilene 2000.    156      780
## 7 Abilene 2000.    152      742
## 8 Abilene 2001.    131      765
## 9 Abilene 2001.    104      771
## 10 Abilene 2001.    101      764
## # ... with 8,592 more rows
```

selecting columns with select()

Use case: You want to present a subset of your columns

```
select(texas_housing_data, -c(city, date, sales, listings))
```

The - says to exclude the columns listed in the vector.

selecting columns with select(), helpers

Use case: You want to reorder your columns

```
select(texas_housing_data, city, date,  
       sales, listings, everything())
```

```
## # A tibble: 8,602 x 9
```

```
##   city      date sales listings  year month  volume med  
##   <chr>   <dbl> <dbl>    <dbl> <int> <int>    <dbl> <dbl>  
## 1 Abilene 2000      72      701  2000     1  5380000  71  
## 2 Abilene 2000.      98      746  2000     2  6505000  58  
## 3 Abilene 2000.     130      784  2000     3  9285000  58  
## 4 Abilene 2000.      98      785  2000     4  9730000  68  
## 5 Abilene 2000.     141      794  2000     5 10590000  67  
## 6 Abilene 2000.     156      780  2000     6 13910000  66  
## 7 Abilene 2000.     152      742  2000     7 12635000  73  
## 8 Abilene 2001.     131      765  2000     8 10710000  75  
## 9 Abilene 2001.     104      771  2000     9  7615000  64  
## 10 Abilene 2001.     101      764  2000    10  7040000  59  
## # ... with 8,592 more rows
```

sort rows with arrange()

storms

arrange()

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



storm	wind	pressure	date
Ana	40	1013	1997-07-01
Alex	45	1009	1998-07-30
Arthur	45	1010	1996-06-21
Arlene	50	1010	1999-06-13
Allison	65	1005	1995-06-04
Alberto	110	1007	2000-08-12

sort rows with arrange()

```
arrange(texas_housing_data, year)
```

```
## # A tibble: 8,602 x 9
```

```
##   city      year month sales    volume median listings in
```

```
##   <chr>    <int> <int> <dbl>    <dbl>    <dbl>    <dbl>
```

```
## 1 Abilene  2000     1     72  5380000  71400     701
```

```
## 2 Abilene  2000     2     98  6505000  58700     746
```

```
## 3 Abilene  2000     3    130  9285000  58100     784
```

```
## 4 Abilene  2000     4     98  9730000  68600     785
```

```
## 5 Abilene  2000     5    141 10590000  67300     794
```

```
## 6 Abilene  2000     6    156 13910000  66900     780
```

```
## 7 Abilene  2000     7    152 12635000  73500     742
```

```
## 8 Abilene  2000     8    131 10710000  75000     765
```

```
## 9 Abilene  2000     9    104  7615000  64500     771
```

```
## 10 Abilene 2000    10    101  7040000  59300     764
```

```
## # ... with 8,592 more rows
```

sort rows with arrange()

To change the order of use desc()

```
arrange(texas_housing_data, desc(year))
```

```
## # A tibble: 8,602 x 9
```

```
##   city      year month sales   volume median listings  
##   <chr>    <int> <int> <dbl>    <dbl>   <dbl>    <dbl>  
## 1 Abilene   2015     1   158 23486998 134100      801  
## 2 Abilene   2015     2   151 19834263 126500      767  
## 3 Abilene   2015     3   198 31869437 136800      821  
## 4 Abilene   2015     4   201 28301159 129600      891  
## 5 Abilene   2015     5   199 31385757 144700      919  
## 6 Abilene   2015     6   260 41396230 141500      965  
## 7 Abilene   2015     7   268 45845730 148700      986  
## 8 Amarillo  2015     1   204 33188726 138500     1120  
## 9 Amarillo  2015     2   188 34355428 149400     1084  
## 10 Amarillo 2015     3   317 53603130 140900     1051  
## # ... with 8,592 more rows
```

Introducing the pipe operator



Interlude: Ceci est une %>%

The pipe %>% operator takes the left-hand side and makes it *input* in the right-hand side.

- ▶ by default, the left-hand side is the *first argument* of the right-hand side function.

```
# a tibble is the first argument
select(texas_housing_data, city, year, sales, volume)

texas_housing_data %>%
  select(city, year, sales, volume)
```

Ceci est une %>%

We can chain together tidyverse functions to avoid making so many intermediate data frames!

```
texas_housing_data %>%  
  select(city, year, month, median) %>%  
  arrange(desc(median))
```

```
## # A tibble: 8,602 x 4  
##       city          year month median  
##   <chr>         <int> <int>   <dbl>  
## 1 Collin County  2015     5 304200  
## 2 Collin County  2015     6 300400  
## 3 Collin County  2015     7 292600  
## 4 Collin County  2015     4 291400  
## 5 Collin County  2015     3 285800  
## 6 Fort Bend     2015     6 284200  
## 7 Collin County  2015     2 283400  
## 8 Midland       2014     6 283100  
## 9 Fort Bend     2014     6 282300
```

creating columns with mutate()

mutate()

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



storm	wind	pressure	date	ratio	inverse
Alberto	110	1007	2000-08-12	9.15	0.11
Alex	45	1009	1998-07-30	22.42	0.04
Allison	65	1005	1995-06-04	15.46	0.06
Ana	40	1013	1997-07-01	25.32	0.04
Arlene	50	1010	1999-06-13	20.20	0.05
Arthur	45	1010	1996-06-21	22.44	0.04

creating columns with mutate()

```
texas_housing_data %>%  
  mutate(mean_price = volume / sales) %>%  
  select(city, year, month, mean_price, sales, volume)
```

```
## # A tibble: 8,602 x 6
```

	city	year	month	mean_price	sales	volume
	<chr>	<int>	<int>	<dbl>	<dbl>	<dbl>
## 1	Abilene	2000	1	74722.	72	5380000
## 2	Abilene	2000	2	66378.	98	6505000
## 3	Abilene	2000	3	71423.	130	9285000
## 4	Abilene	2000	4	99286.	98	9730000
## 5	Abilene	2000	5	75106.	141	10590000
## 6	Abilene	2000	6	89167.	156	13910000
## 7	Abilene	2000	7	83125	152	12635000
## 8	Abilene	2000	8	81756.	131	10710000
## 9	Abilene	2000	9	73221.	104	7615000
## 10	Abilene	2000	10	69703.	101	7040000

```
## # ... with 8,592 more rows
```

Binary operators: Math in R

R is a calculator! We can do math with numbers, using the following symbols:

4 + 4

4 - 4

4 * 4

4 / 4

4 ^ 4

5 %% 4 *# gives the remainder after dividing*

creating columns with mutate()

When we mutate, you can create new columns.

- ▶ On the right side of the equal sign, you have the name of a new column.
- ▶ On the left side, you have code that creates a new column (using vector operations)¹

```
texas_housing_data %>%  
  mutate(mean_price = volume / sales) %>%  
  select(city, year, month, mean_price, sales, volume)
```

```
## # A tibble: 8,602 x 6  
##   city      year month mean_price sales  volume  
##   <chr>   <int> <int>      <dbl> <dbl>   <dbl>  
## 1 Abilene  2000     1    74722.    72 5380000  
## 2 Abilene  2000     2    66378.    98 6505000  
## 3 Abilene  2000     3    71423.   130 9285000  
## 4 Abilene  2000     4    99286.    98 9730000  
## 5 Abilene  2000     5    75106.   141 10590000  
## 6 Abilene  2000     6    88167.   156 13810000
```

creating columns with mutate()

You can create multiple columns at a single time and even use information from a newly created column as input.

```
texas_housing_data %>%  
  mutate(mean_price = volume / sales,  
         sqrt_mean_price = sqrt(mean_price)) %>%  
  select(city, year, month, mean_price, sales, volume)
```

```
## # A tibble: 8,602 x 6
```

```
##   city      year month mean_price sales  volume  
##   <chr>    <int> <int>      <dbl> <dbl>   <dbl>  
## 1 Abilene  2000     1    74722.    72 5380000  
## 2 Abilene  2000     2    66378.    98 6505000  
## 3 Abilene  2000     3    71423.   130 9285000  
## 4 Abilene  2000     4    99286.    98 9730000  
## 5 Abilene  2000     5    75106.   141 10590000  
## 6 Abilene  2000     6    89167.   156 13910000  
## 7 Abilene  2000     7    83125    152 12635000  
## 8 Abilene  2000     8    81756.   131 10710000
```

choose rows that match a condition with `filter()`

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



`filter()`

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Ana	40	1013	1997-07-01

choose rows that match a condition with filter()

Get all the data from 2013

```
filter(texas_housing_data, year == 2013)
```

```
## # A tibble: 552 x 9
```

```
##   city      year month sales    volume median listings in
```

```
##   <chr>    <int> <int> <dbl>    <dbl>  <dbl>    <dbl>
```

```
## 1 Abilene  2013     1   114 15794494 125300     966
```

```
## 2 Abilene  2013     2   140 16552641  94400     943
```

```
## 3 Abilene  2013     3   164 19609711 102500     958
```

```
## 4 Abilene  2013     4   213 27261796 113700     948
```

```
## 5 Abilene  2013     5   225 31901380 130000     923
```

```
## 6 Abilene  2013     6   209 29454125 127300     960
```

```
## 7 Abilene  2013     7   218 32547446 140000     969
```

```
## 8 Abilene  2013     8   236 30777727 120000     976
```

```
## 9 Abilene  2013     9   195 26237106 127500     985
```

```
## 10 Abilene 2013    10   167 21781187 119000     993
```

```
## # ... with 542 more rows
```

Relational operators return TRUE or FALSE

Before moving forward with `filter()`, we need to know about relational operators and logical operators

Operator	Name
<	less than
>	greater than
<=	less than or equal to
>=	greater than or equal to
==	equal to
!=	not equal to
%in%	matches something in

Relational operators in practice

```
4 < 4
```

```
## [1] FALSE
```

```
4 >= 4
```

```
## [1] TRUE
```

```
4 == 4
```

```
## [1] TRUE
```

```
4 != 4
```

```
## [1] FALSE
```

```
4 %in% c(1, 2, 3)
```

```
## [1] FALSE
```

logical operators combine TRUEs and FALSEs logically

Operator	Name
!	not
&	and
	or

```
# not true
```

```
! TRUE
```

```
## [1] FALSE
```

```
# are both x & y TRUE?
```

```
TRUE & FALSE
```

```
## [1] FALSE
```

```
# is either x | y TRUE?
```

```
TRUE | FALSE
```

```
## [1] TRUE
```

What do the following return?

Logical operators team up with relational operators.

- ▶ First, evaluate the relational operator
- ▶ Then, carry out the logic.

```
! (4 > 3) # ! TRUE
```

```
(5 > 1) & (5 > 2) # TRUE & TRUE
```

```
(4 > 10) | (20 > 3) # FALSE | TRUE
```

This is hard to wrap your head around. We'll have plenty of practice!

choose rows that match a condition with filter()

Get all the data from 2013 for Houston.

- ▶ in filter() additional match criteria are treated like and

```
texas_housing_data %>%  
  filter(year == 2013,  
         city == "Houston")
```

```
## # A tibble: 12 x 9
```

##	city	year	month	sales	volume	median	listings
##	<chr>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 Houston	2013	1	4273	852045057	149500	21364
##	2 Houston	2013	2	4886	1060985674	161900	21293
##	3 Houston	2013	3	6382	1479273481	172300	20909
##	4 Houston	2013	4	7116	1770746764	182400	20607
##	5 Houston	2013	5	8439	2121508529	186100	20526
##	6 Houston	2013	6	7935	2073909387	191600	21008
##	7 Houston	2013	7	8468	2168720825	187800	21497
##	8 Houston	2013	8	8155	2083377894	186700	21366
##	9 Houston	2013	9	6706	1628002370	180000	21007

choose rows that match a condition with `filter()`

Get all the data from 2013 for Houston or Austin

- ▶ in `filter()` additional match criteria are treated like and
- ▶ we get nothing returned here, because no observation is in Houston AND in Austin.

```
texas_housing_data %>%  
  filter(year == 2013,  
         city == "Houston", city == "Austin")
```

```
## # A tibble: 0 x 9
```

```
## # ... with 9 variables: city <chr>, year <int>, month <int>
```

```
## #   volume <dbl>, median <dbl>, listings <dbl>, inventory <dbl>
```

choose rows that match a condition with filter()

Get all the data from after than 2013 for Houston OR Austin

```
texas_housing_data %>%  
  filter(year > 2013,  
         city == "Houston" | city == "Austin")
```

```
## # A tibble: 38 x 9
```

##	city	year	month	sales	volume	median	listings
##	<chr>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 Austin	2014	1	1582	426127544	213700	5118
##	2 Austin	2014	2	1903	550882376	229400	5255
##	3 Austin	2014	3	2434	717821612	235600	5512
##	4 Austin	2014	4	2691	813253968	237000	5838
##	5 Austin	2014	5	3178	1012123948	243900	6539
##	6 Austin	2014	6	3195	1023051880	248900	7040
##	7 Austin	2014	7	3151	982086356	246900	7475
##	8 Austin	2014	8	3023	927019222	243800	7326
##	9 Austin	2014	9	2664	813797562	238900	7072
##	10 Austin	2014	10	2588	796863816	239600	6791

choose rows that match a condition with filter()

Get all the data from after than 2013 for Houston Galveston

```
texas_housing_data %>%  
  filter(year > 2013,  
         city %in% c("Houston", "Dallas", "Austin"))
```

```
## # A tibble: 57 x 9
```

##	city	year	month	sales	volume	median	listings
##	<chr>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 Austin	2014	1	1582	426127544	213700	5118
##	2 Austin	2014	2	1903	550882376	229400	5255
##	3 Austin	2014	3	2434	717821612	235600	5512
##	4 Austin	2014	4	2691	813253968	237000	5838
##	5 Austin	2014	5	3178	1012123948	243900	6539
##	6 Austin	2014	6	3195	1023051880	248900	7040
##	7 Austin	2014	7	3151	982086356	246900	7475
##	8 Austin	2014	8	3023	927019222	243800	7326
##	9 Austin	2014	9	2664	813797562	238900	7072
##	10 Austin	2014	10	2588	796863816	239600	6791

summarize data with `summarize()`

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



median
22.5

summarize data with summarize()

Calculate total volume of sales in Texas from 2014.

```
texas_housing_data %>%  
  filter(year == 2014) %>%  
  summarize(total_volume = sum(volume))
```

```
## # A tibble: 1 x 1  
##   total_volume  
##   <dbl>  
## 1 84760948831
```

summarize data with summarize()

Calculate the mean and median number of sales in Texas's three largest cities.

```
texas_housing_data %>%  
  filter(city %in%  
          c("Houston", "Dallas", "San Antonio")) %>%  
  summarize(median_n_sales = median(sales),  
            mean_n_sales = mean(sales))
```

```
## # A tibble: 1 x 2  
##   median_n_sales mean_n_sales  
##           <dbl>         <dbl>  
## 1           3996           3890.
```

summarize data with summarize()

There are many useful functions that go with summarize. Try `?summarize` for more.

```
texas_housing_data %>%  
  filter(city %in%  
           c("Houston", "Dallas", "San Antonio")) %>%  
  summarize(n_obs = n(),  
            n_cities = n_distinct(city))
```

```
## # A tibble: 1 x 2  
##   n_obs n_cities  
##   <int>   <int>  
## 1    561       3
```


summarize data with summarize()

If you try to make a summarize statistic that does not collapse the data to a single value (per group), you'll get an error like so:

```
texas_housing_data %>%  
  filter(city %in%  
           c("Houston", "Dallas", "San Antonio")) %>%  
  summarize(mean_price = volume / sales)
```

Error: Column `mean_price` must be length 1 (a summary value)

Get number of observations

piping dplyr verbs together

dplyr verbs can be piped together in any order you want, although different orders can give you different results, so be careful!

```
texas_housing_data %>%  
  select(city, year, month, sales, volume) %>%  
  mutate(log_mean_price = log(volume / sales)) %>%  
  filter(year == 2013) %>%  
  summarize(log_mean_price_2013 = mean(log_mean_price, na.rm = TRUE))
```

```
## # A tibble: 1 x 1  
##   log_mean_price_2013  
##               <dbl>  
## 1                12.1
```

```
# Won't give you the same result as  
# texas_housing_data %>%  
#   select(city, year, month, sales, volume) %>%  
#   mutate(log_mean_price = log(volume / sales)) %>%  
#   summarize(log_mean_price = mean(log_mean_price, na.rm = TRUE))  
#   filter(year == 2013)
```

Recap: manipulating data with dplyr

We learned

- ▶ how to employ the 5 dplyr verbs of highest importance including
 - ▶ `select()` to pick columns
 - ▶ `arrange()` to order the data
 - ▶ `mutate()` to create new columns
 - ▶ `filter()` to get rows that meet a criteria
 - ▶ `summarize()` to summarize data
- ▶ how to use relation operators, binary operators for math and logical operators in dplyr contexts