

Coding Lab: Visualizing data with ggplot2

Ari Anisfeld

Summer 2020

How to use ggplot

- ▶ How to map data to aesthetics with `aes()` (and what that means)
- ▶ How to visualize the mappings with `geoms`
- ▶ How to get more out of your data by using multiple aesthetics
- ▶ How to use facets to add dimensionality

There are whole books on how to use ggplot. This is a quick introduction!

Understanding ggplot()

By itself, ggplot() tells R to prepare to make a plot.


```
texas_annual_sales <-  
  texas_housing_data %>%  
  group_by(year) %>%  
  summarize(total_volume = sum(volume, na.rm = TRUE))  
  
ggplot(data = texas_annual_sales)
```

Adding a mapping

Adding `mapping = aes()` says how the data will map to “aesthetics”.

- ▶ e.g. tell R to make x-axis year and y-axis `total_volume`.
- ▶ Each row of the data has (year, `total_volume`).
 - ▶ R will map that to the coordinate pair (x,y) .
 - ▶ Look at the data before moving on!

```
ggplot(data = texas_annual_sales,  
       mapping = aes(x = year, y = total_volume))
```



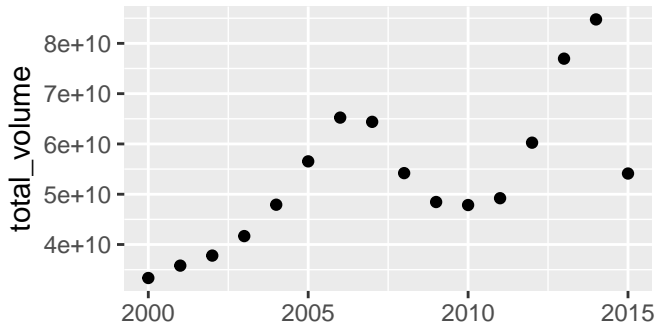
Visualizing the mapping with a geom

`geom_<name>` tells R what type of visualization to produce.

Here we see points.

- ▶ Each row of the data has (year, total_volume).
- ▶ R will map that to the coordinate pair (x,y).

```
ggplot(data = texas_annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_point()
```




Visualizing the mapping with a geom

Here we see bars.

- ▶ Each row of the data has (year, total_volume).
- ▶ R will map that to the coordinate pair (x,y)


```
ggplot(data = texas_annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_col()
```



Visualizing the mapping with a geom

Here we see a line connecting each (x,y) pair.

```
ggplot(data = texas_annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_line()
```




Visualizing the mapping with a geom

Here we see a smooth line. R does a statistical transformation!

- ▶ Now R doesn't visualize the mapping (year, total_volume) to each (x,y) pair
- ▶ Instead it fits a model to the (x,y) and then plots the “smooth” line

```
ggplot(data = texas_annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_smooth(se = FALSE)
```


`geom_smooth()` using method = 'loess' and formula 'y ~ x'



Visualizing the mapping with a geom

We can overlay several geom.

```
ggplot(data = texas_annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_smooth(se = FALSE) +  
  geom_point()
```



Visualizing the mapping with a geom


- ▶ We saw that we can visualize a relationship between two variables mapping data to x and y
- ▶ The data can be visualized with different geoms that can be composed (+) together.
- ▶ We can even calculate new variables with statistics and plot those on the fly.

Next: Now we'll look at aesthetics that go beyond x and y axes.

Using aesthetics to explore data.

We'll use midwest data and start with only mapping to x and y


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty)) +  
  geom_point()
```



Using aesthetics to explore data.

- ▶ color maps data to the color of points or lines.
 - ▶ Each state is assigned a color.
 - ▶ This works with discrete data and continuous data.


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             color = state)) +  
  geom_point()
```



Using aesthetics to explore data.

- ▶ shape maps data to the shape of points.
 - ▶ Each state is assigned a shape.
 - ▶ This works with discrete data only.

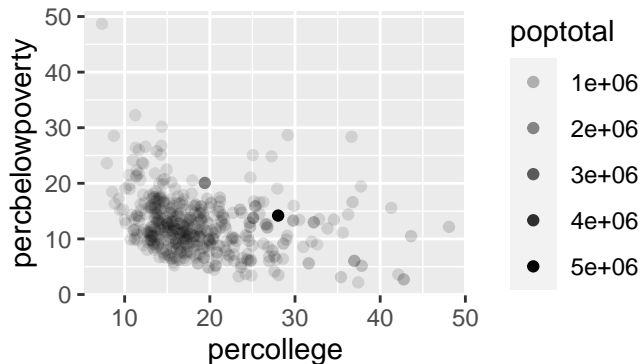
```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             shape = state)) +  
  geom_point()
```



Using aesthetics to explore data.

- ▶ alpha maps data to the transparency of points.
 - ▶ Here we map the percentage of people within a known poverty status to alpha


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             alpha = poptotal)) +  
  geom_point()
```



Using aesthetics to explore data.

- ▶ size maps data to the size of points and width of lines.
- ▶ Here we map the percentage of people within a known poverty status to size

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             size = poptotal)) +  
  geom_point()
```

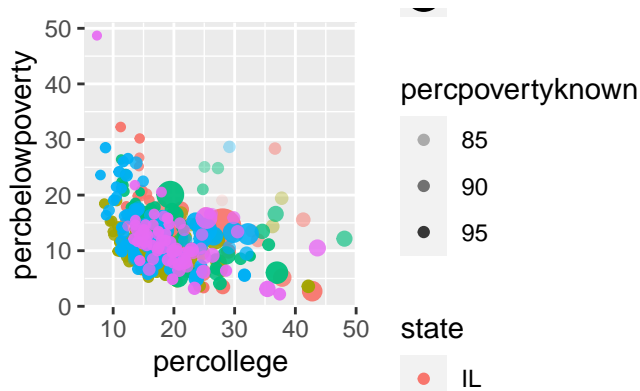


Using aesthetics to explore data.

We can combine any and all aesthetics, and even map the same variable to multiple aesthetics

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             alpha = percpovertyknown,  
             size = poptotal,  
             color = state)) +  
  geom_point()
```


Using aesthetics to explore data.



Using aesthetics to explore data


Different geoms have specific aesthetics that go with them.

- ▶ use `?` to see which aesthetics a geom accepts (e.g `?geom_point`)
 - ▶ the bold aesthetics are required.
- ▶ the ggplot cheatsheet shows all the geoms with their associated aesthetics

Facets

Facets provide an additional tool to explore multidimensional data

```
midwest %>%  
  ggplot(aes(x = log(poptotal),  
             y = percbelowpoverty)) +  
  geom_point() +  
  facet_wrap(vars(state))
```




discrete

vs continuous data

color can be continuous

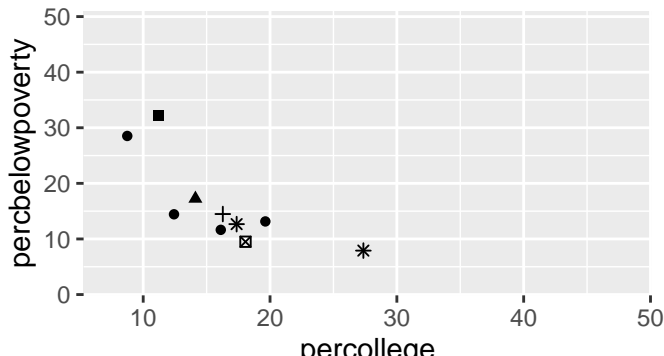
```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             color = percpovertyknown)) +  
  geom_point()
```



shape does not play well with many categories

- ▶ Will only map to 6 categories, the rest become NA.
- ▶ We can override this behavior and get up to 25 distinct shapes


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             shape = county)) +  
  geom_point() +  
  # legend off, otherwise it overwhelms  
  theme(legend.position = "none")
```



alpha and size can be misleading with discrete data


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             alpha = state)) +  
  geom_point()
```

Warning: Using alpha for a discrete variable is not advised



Adding vertical lines

```
texas_annual_sales %>%  
  ggplot(aes(x = year, y = total_volume)) +  
    geom_point() +  
    geom_vline(aes(xintercept = 2007),  
               linetype = "dotted")
```



- ▶ add horizontal lines with `geom_hline()`
- ▶ add any linear fit using `geom_abline()` by providing a slope

Key take aways

- ▶ `ggplot` starts by mapping data to “aesthetics”.
 - ▶ e.g. What data shows up on x and y axes and how color, size and shape appear on the plot.
 - ▶ We need to be aware of ‘continuous’ vs. ‘discrete’ variables.
- ▶ Then, we use `geoms` to create a visualization based on the mapping.
 - ▶ Again we need to be aware of ‘continuous’ vs. ‘discrete’ variables.
- ▶ Making quick plots helps us understand data and makes us aware of data issues

Resources: R for Data Science chap. 3 (r4ds.had.co.nz); RStudio's `ggplot` cheatsheet.


Appendix: Some graphs you made along the way

lab 0: a map

`geom_path` is like `geom_line`, but connects (x, y) pairs in the order they appear in the data set.

```
storms %>%  
  group_by(name, year) %>%  
  filter(max(category) == 5) %>%  
ggplot(aes(x = long, y = lat, color = name)) +  
  geom_path() +  
  borders("world") +  
  coord_quickmap(xlim = c(-130, -60), ylim = c(20, 50))
```


lab 0: a map



lab 1: a line plot

```
french_data <-  
  wid_data %>%  
  filter(type == "Net personal wealth",  
         country == "France") %>%  
  mutate(perc_national_wealth = value * 100)  
  
french_data %>%  
  ggplot(aes(y = perc_national_wealth,  
            x = year,  
            color = percentile)) +  
  geom_line()
```

lab 1: a line plot




lab 2: distributions

- ▶ `geom_density()` only requires an `x` aesthetic and it calculates the distribution to plot.
- ▶ We can set the aesthetics manually, independent of data for nicer graphs.

```
chi_sq_samples <-  
  tibble(x = c(rchisq(100000, 2),  
               rchisq(100000, 3),  
               rchisq(100000, 4)),  
         df = rep(c("2", "3", "4"), each = 1e5))
```

```
chi_sq_samples %>%  
  ggplot(aes(x = x, fill = df)) +  
  geom_density(alpha = .5) +  
  labs(fill = "df", x = "sample")
```

lab 2: distributions




lab 4: grouped bar graphs

- ▶ `position = "dodge2"` tells R to put bars next to each other, rather than stacked on top of each other.
- ▶ Notice we use `fill` and not `color` because we're "filling" an area.

```
mean_share_per_country %>%  
  ggplot(aes(y = country,  
             x = mean_share,  
             fill = percentile)) +  
  geom_col(position = "dodge2") +  
  labs(x = "Mean share of national wealth",  
       y = "",  
       fill = "Wealth\npercentile")
```


lab 4: grouped bar graphs



lab 4: faceted bar graph

- ▶ Notice that we manipulate our data to the right specification before making this graph
- ▶ Using `facet_wrap` we get a distinct graph for each time period.

```
mean_share_per_country_with_time %>%  
  ggplot(aes(x = country,  
             y = mean_share,  
             fill = percentile)) +  
  geom_col(position = "dodge2") +  
  facet_wrap(vars(time_period))
```

lab 4: faceted bar graph

