

# Coding Lab: Functions

Ari Anisfeld

Fall 2020

# Functions

```
# example of a function  
circle_area <- function(r) {  
  
  pi * r ^ 2  
  
}
```

- ▶ What are functions and why do we want to use them?
- ▶ How do we write functions in practice?
- ▶ What are some solutions to avoid frustrating code?

# Motivation

*“You should consider writing a function whenever you’ve copied and pasted a block of code more than twice (i.e. you now have three copies of the same code)”*

- ▶ *Hadley Wickham, R for Data Science*

## Instead of repeating code . . .

```
data %>%  
  mutate(a = (a - min(a)) / (max(a) - min(a)),  
         b = (b - min(b)) / (max(b) - min(b)),  
         c = (c - min(c)) / (max(c) - min(c)),  
         d = (d - min(d)) / (max(d) - min(d)))
```

```
## # A tibble: 100 x 4  
##       a      b      c      d  
##   <dbl> <dbl> <dbl> <dbl>  
## 1 0.0924 0.155 0.531 0.172  
## 2 0.732  0.400 0.688 0.835  
## 3 0.493  0.797 0.146 0.479  
## 4 0.478  0.556 0.631 0.244  
## 5 0.514  0.423 0.523 0.534  
## 6 0.681  0.766 0.295 0.663  
## 7 0.835  0.366 0.840 0.225  
## 8 0.370  0.180 0.607 0.462  
## 9 0.227  0.636 0.938 0.237  
## 10 0.845  0.488 0.981 0.626
```

## Write a function

```
rescale_01 <- function(x) {  
  (x - min(x)) / (max(x) - min(x))  
}
```

```
data %>%  
  mutate(a = rescale_01(a),  
         b = rescale_01(b),  
         c = rescale_01(c),  
         d = rescale_01(d))
```

```
## # A tibble: 100 x 4  
##       a      b      c      d  
##   <dbl> <dbl> <dbl> <dbl>  
## 1 0.0924 0.155 0.531 0.172  
## 2 0.732  0.400 0.688 0.835  
## 3 0.493  0.797 0.146 0.479  
## 4 0.478  0.556 0.631 0.244  
## 5 0.514  0.422 0.522 0.524
```

# Function anatomy

The anatomy of a function is as follows:

```
function_name <- function(arguments) {  
  do_this(arguments)  
}
```

A function consists of

1. Function arguments<sup>1</sup>
2. Function body

We can assign the function to a name like any other object in R.

---

<sup>1</sup>Tech detail: R refers to these as formals.

## Function anatomy: example

- ▶ **arguments:** x
- ▶ **body:** (x - min(x)) / (max(x) - min(x))
- ▶ assign to **name:** rescale\_01

```
rescale_01 <- function(x) {  
  (x - min(x)) / (max(x) - min(x))  
}
```

Note that we don't need to explicitly call `return()`

- ▶ the last line of the code will be the value returned by the function.

## Writing a function: printing output

You start writing code to say Hello to all of your friends.

- ▶ You notice it's getting repetitive. ... time for a function

```
print("Hello Jasmin!")
```

```
## [1] "Hello Jasmin!"
```

```
print("Hello Joan!")
```

```
## [1] "Hello Joan!"
```

```
print("Hello Andrew!")
```

```
## [1] "Hello Andrew!"
```

```
# and so on...
```



# Writing a function: parameterize the code

Start with the **body**.

Ask: What part of the code is changing?

- ▶ Make this an **argument**

## Writing a function: parameterize the code

Start with the **body**.

Rewrite the code to accommodate the parameterization

```
# print("Hello Jasmin!") becomes ...
```

```
name <- "Jasmin"
```

```
print(paste0("Hello ", name, "!"))
```

```
## [1] "Hello Jasmin!"
```

Check several potential inputs to avoid future headaches

## Writing a function: add the structure

```
# name <- "Jasmin"  
# print(paste0("Hello ", name, "!"))  
  
function(name) {  
  print(paste0("Hello ", name, "!"))  
}
```

## Writing a function: assign to a name

Try to use **names** that actively tell the user what the code does

- ▶ We recommend `verb_thing()`
  - ▶ **good** `calc_size()` or `compare_prices()`
  - ▶ **bad** `prices()`, `calc()`, or `fun1()`.

```
# name <- "Jasmin"
# print(paste0("Hello ", name, "!"))

say_hello_to <- function(name) {
  print(paste0("Hello ", name, "!"))
}
```

## Simple example: printing output

Test out different inputs!

```
say_hello_to("Jasmin")
```

```
## [1] "Hello Jasmin!"
```

```
say_hello_to("Joan")
```

```
## [1] "Hello Joan!"
```

```
say_hello_to(name = "Andrew")
```

```
## [1] "Hello Andrew!"
```

```
# Cool this function is vectorized!
```

```
say_hello_to(c("Jasmin", "Joan", "Andrew"))
```

```
## [1] "Hello Jasmin!" "Hello Joan!"   "Hello Andrew!"
```

Question: does name exist in my R environment after I run this function? Why or why not?

## Technical aside: `typeof(your_function)`

Like other R objects functions have types.

Primitive functions are of type “builtin”

```
typeof(`+`)
```

```
## [1] "builtin"
```

```
typeof(sum)
```

```
## [1] "builtin"
```

## Technical aside: `typeof(your_function)`

Like other R objects functions have types.

User defined functions, functions loaded with packages and many base R functions are type “closure”:

```
typeof(say_hello_to)
```

```
## [1] "closure"
```

```
typeof(mean)
```

```
## [1] "closure"
```

## Technical aside: `typeof(your_function)`

This is background knowledge that might help you understand an error.

For example, you thought you assigned a number to the name “c” and want to calculate ratio.

```
ratio <- 1 / c
```

```
Error in 1/c : non-numeric argument to binary operator  
as.integer(c)
```

Error in `as.integer(c)` :

cannot coerce type 'builtin' to vector of type 'integer'

“builtin” or “closure” in this situation let you know your working with a function!



## Second example: calculating the mean of a sample

Your stats prof asks you to simulate a central limit theorem, by calculating the mean of samples from the standard normal distribution with increasing sample sizes.

```
mean(rnorm(1))
```

```
## [1] 0.09486633
```

```
mean(rnorm(3))
```

```
## [1] -0.3880646
```

```
mean(rnorm(30))
```

```
## [1] 0.1469901
```

```
# et cetera
```

## Second example: calculating the mean of a sample

The number is changing, so it becomes the **argument**.

```
calc_sample_mean <- function(sample_size) {  
  
  mean(rnorm(sample_size))  
  
}
```

- ▶ The number is the sample size, so I call it `sample_size`. `n` would also be appropriate.
- ▶ The **body** code is otherwise identical to the code you already wrote.

## Second example: calculating the mean of a sample

For added clarity you can unnest your code and assign the intermediate results to meaningful names.

```
calc_sample_mean <- function(sample_size) {  
  
  random_sample <- rnorm(sample_size)  
  
  sample_mean <- mean(random_sample)  
  
  return(sample_mean)  
}
```

`return()` explicitly tells R what the function will return.

- ▶ The last line of code run is returned by default.

## Second example: calculating the mean of a sample

If the function can be fit on one line, then you can write it without the curly brackets like so:

```
calc_sample_mean <- function(n) mean(rnorm(n))
```

Some settings call for anonymous functions, where the function has no name.

```
function(n) mean(rnorm(n))
```

```
## function(n) mean(rnorm(n))
```

# Always test your code

Try to foresee the kind of input you expect to use.

```
calc_sample_mean(1)
```

```
## [1] -0.7721886
```

```
calc_sample_mean(1000)
```

```
## [1] -0.007513187
```

We see below that this function is not vectorized. We might hope to get 3 sample means out but only get 1

```
# read ?rnorm to understand how rnorm  
# interprets vector input.  
calc_sample_mean(c(1, 3, 30))
```

```
## [1] -0.3832335
```

## Adding additional arguments

If we want to be able to adjust the details of how our function runs we can add arguments

- ▶ typically, we put “data” arguments first
- ▶ and then “detail” arguments after

```
calc_sample_mean <- function(sample_size,  
                               our_mean,  
                               our_sd) {  
  
  sample <- rnorm(sample_size,  
                  mean = our_mean,  
                  sd = our_sd)  
  
  mean(sample)  
}
```

## Setting defaults

We usually set default values for “detail” arguments.

```
calc_sample_mean <- function(sample_size,
                                our_mean = 0,
                                our_sd = 1) {

  sample <- rnorm(sample_size,
                  mean = our_mean,
                  sd = our_sd)

  mean(sample)
}
```

```
# uses the defaults
```

```
calc_sample_mean(sample_size = 10)
```

```
## [1] -0.2030251
```

## Setting defaults

```
# we can change one or two defaults.  
# You can refer by name, or use position  
calc_sample_mean(10, our_sd = 2)
```

```
## [1] -0.0548051
```

```
calc_sample_mean(10, our_mean = 6)
```

```
## [1] 6.348498
```

```
calc_sample_mean(10, 6, 2)
```

```
## [1] 5.870006
```



## Setting defaults

This won't work though:

```
calc_sample_mean(our_mean = 5)
```

```
Error in rnorm(sample_size, mean = our_mean, sd = our_sd) :  
  argument "sample_size" is missing, with no default
```

## Key points

- ▶ Write functions when you are using a set of operations repeatedly
- ▶ Functions consist of arguments and a body and are usually assigned to a name.
- ▶ Functions are for humans
  - ▶ pick names for the function and arguments that are clear and consistent
- ▶ Debug your code as much as you can as you write it.
  - ▶ if you want to use your code with `mutate()` test the code with vectors

**For more:** See Functions Chapter in R for Data Science

Additional material

## Functions in functions

We can pass functions as arguments to other functions. Before:

```
calc_sample_mean <- function(sample_size,
                              our_mean = 0,
                              our_sd = 1) {
  sample_mean <- mean(rnorm(sample_size,
                             mean = our_mean,
                             sd = our_sd))

  sample_mean
}
```

## Functions in functions

We can pass functions as arguments to other functions. After:

```
summarize_sample <- function(sample_size,
                              our_mean = 0,
                              our_sd = 1,
                              summary_fxn = mean) {
  summary_stat <- summary_fxn(rnorm(sample_size,
                                     mean = our_mean,
                                     sd = our_sd))

  summary_stat
}
```

## Functions in functions

```
calc_sample_mean(sample_size = 10,  
                  our_mean = 0,  
                  our_sd = 1)
```

```
## [1] -0.2553635
```

```
summarize_sample(sample_size = 10,  
                  our_mean = 0,  
                  our_sd = 1,  
                  summary_fxn = max)
```

```
## [1] 1.744373
```

`calc_sample_mean()` is now probably the wrong name for this function - we should call it `summarize_sample()` or something like that.

## Detour: probability distributions

R has built-in functions for working with distributions.

	example	what it does?
d	<code>dnorm(x)</code>	returns pdf value at $x$
p	<code>pnorm(q)</code>	returns CDF value at $q$
q	<code>qnorm(p)</code>	returns inverse CDF (the quantile) for a given probability
r	<code>rnorm(n)</code>	generates a random sample of size $n$

Probability distributions you are familiar with are likely built-in to R.

For example, the binomial distribution has `dbinom()`, `pbinom()`, `qbinom()`, `rbinom()`. The  $t$  distribution has `dt()`, `pt()`, `qt()`, `rt()`, etc.

Read this tutorial for more information.

## Detour: probability distributions

- ▶ `dnorm()`: density function, the PDF evaluated at  $X$ .

```
dnorm(0)
```

```
## [1] 0.3989423
```

```
dnorm(1)
```

```
## [1] 0.2419707
```

```
dnorm(-1)
```

```
## [1] 0.2419707
```



## Detour: probability distributions

## Detour: probability distributions

- ▶ `pnorm()`: cumulative distribution function, the CDF evaluated at  $X$ .

```
pnorm(0)
```

```
## [1] 0.5
```

```
pnorm(1)
```

```
## [1] 0.8413447
```

```
pnorm(-1)
```

```
## [1] 0.1586553
```

## Detour: probability distributions

## Detour: probability distributions

- `qnorm()`: quantile function, the inverse CDF evaluated at a quantile.

```
qnorm(c(0.05, 0.95))
```

```
## [1] -1.644854  1.644854
```

```
qnorm(c(0.025, 0.975))
```

```
## [1] -1.959964  1.959964
```

```
pnorm(qnorm(c(0.025, 0.975)))
```

```
## [1] 0.025 0.975
```

## Detour: probability distributions

- `rnorm()`: random sampling

```
rnorm(1)
```

```
## [1] 0.7244536
```

```
rnorm(5)
```

```
## [1] -1.4225062 -0.5960598  0.3353115  0.1752422 -1.24801
```

```
rnorm(30)
```

```
## [1] 0.7412269020  0.2417359429 -0.1755121315 -0.269961
```

```
## [6] -0.3039516771 -0.7319315023 -0.9072302122 -0.107411
```

```
## [11] 0.6332928702  0.0033477103 -0.5941185630  0.449093
```

```
## [16] 0.0151888586 -0.9255065013  0.4645399559  1.017972
```

```
## [21] 0.3655853978  0.9052002039 -3.0692947194 -1.129443
```

```
## [26] 0.8857969353  0.0054698593  0.4577974162  0.897436
```